TRAN NHAT NAM

Product Data Engineer | Ho Chi Minh City, Vietnam

+84 832255586 | nam.trannhat067@gmail.com | linkedin.com/in/nam-tran-nhat | github.com/namtran067

Summary

Product Data Engineer with 4+ years of experience creating **end-to-end data solutions** – from **ETL/ELT pipelines** and **data platforms** to **data governance, BI and AI applications**; now operating **2,000+ table ingestion (400+ sources, 100M+ rows/day)** on Databricks at Nhà Thuốc Long Châu. More than a technician: I dive deep into the business needs of each department (Sales, Marketing, Finance, BOD) to turn data into decisions – **cutting ETL operational cost by 90%**, **consultation time by 60%**, and **accelerating decisions up to 3x**.

Technical Skills

- **Languages:** Python, SQL
- **Data Engineering:** Databricks (Delta Lake, medallion architecture), Azure Data Factory, Apache Kafka, dbt (medallion modeling, tests, docs), Dagster, Airflow, dlt, Airbyte, n8n, ETL/ELT
- **Databases & Warehouses:** PostgreSQL, BigQuery, DuckDB, StarRocks, dimensional modeling (star schema), data lineage, vector databases
- **Lakehouse & Processing:** Apache Iceberg, Apache Spark (PySpark), Trino, Polaris, MinIO, Parquet
- **BI & AI/ML:** Looker, Metabase, LLM APIs, RAG & vector search, prompt engineering, HuggingFace, ASR (Wav2Vec 2.0)
- **Cloud, DevOps & Deployment:** Azure, AWS Lambda, Google Kubernetes Engine (GKE), Helm, Docker, GitHub Actions, CI/CD, Linux, Vercel, Railway, Dokploy, Coolify
- **Integration & Business Systems:** REST APIs, web scraping (Scrapy, BeautifulSoup), Odoo, ERPNext, Haravan, BizFly, AppSheet
- **Stakeholder Communication:** translating business requirements into technical specs, executive presentations, self-serve dashboards

Experience

Data Engineer | Nhà Thuốc Long Châu **Aug 2026 – Present**

- **[Ingestion at Scale]** Operate and maintain ingestion of **2,000+ tables** from **400+ source databases** into the bronze (raw) layer – largest tables **100M+ rows/day** – the reliable raw supply every downstream team builds on.
- **[Silver Layer]** Build base transformations from bronze to silver, delivering cleansed, conformed datasets consumed by downstream analytics and data teams.
- **[Platform Migration]** Key contributor migrating pipelines from Azure Data Factory + Databricks to a **full-Databricks platform** – removing the ADF hop to cut latency, cost and operational complexity.
- **[Streaming & Batch Standards]** Stand up streaming (**120+ Kafka topics**) and batch pipelines on Databricks and **define the team's standards** for both patterns – consistent development, monitoring and on-call.

Product Data Engineer | Jemmia Diamond **Sep 2023 – Jul 2026**

- **[Data Platform & Orchestration]** Architected and deployed a Modern Data Stack on Dagster (via Dokploy) as software-defined assets, with dlt/Airbyte/n8n ingestion centralizing 12+ sources (Social Ads, Haravan, ERP, Stringee) into a PostgreSQL Data Warehouse – 150 raw tables processing 500K+ rows/day, with CI/CD and Sentry anomaly alerting throughout.
- **[Data Modeling & Transformation]** Designed Star Schema (Kimball) models and authored 200+ dbt transformations across Raw → Staging → Intermediate → Marts, cutting development and maintenance time by 90% with dbt tests (schema, not-null, relationships, freshness), PII masking and lineage via OpenMetadata.
- **[Serverless ETL & IoT Integration]** Migrated brittle legacy workflows (Zapier) to event-driven architectures on AWS Lambda (deployed via AWS SAM). Designed a high-throughput pipeline integrating hardware RFID scanners with e-commerce systems via webhooks (~1,000 records / 3MB per scan, 20 scans/day), enabling time-series analysis for inventory optimization and supply-chain management.
- **[Enterprise Systems & Master Data]** Led enterprise systems transformation: migrated 10GB+ to ERPNext and engineered a custom PIM on NocoDB (40+ tables) over a Cloudflare R2 data lake, with zero-downtime backfill for 8,000+ SKUs and CDC – cutting cross-departmental manual operations 40% at 97% adoption.
- **[Business Intelligence & Analytics]** Primary analytics partner for Sales, Marketing, Finance and the BOD – defined KPIs and a governed semantic layer, shipped 30+ executive dashboards in Looker (100% key-metric coverage), then migrated BI to Metabase for self-service analytics – accelerating data-driven decisions up to 3x.

- **[GenAI & Data Activation]** Built AI products on the warehouse: **Salesaya** (AI sales assistant cutting consultation time 60%) and **Metabot** – a Telegram/Discord agent for secure natural-language data queries without exposing the underlying database (Metabase API, Next.js, React, Redis context caching).

Data Analyst (Contract) | BabyBet (Startup) **Jul 2023 – Mar 2024**

- **Contracted to assess a mobility venture's feasibility:** consolidated market & competitor data on 5+ players (Grab, be) from 15+ market-intelligence, survey and secondary-research sources into structured datasets.
- **Delivered 10+ demand, willingness-to-pay, pricing and unit-economics analyses** that informed the cross-functional go/no-go decision and phased strategy.

Analytics Engineer | AVLI Group JSC **Sep 2022 – Jul 2023**

- **Built the ingestion layer:** AWS Lambda scraping workers (Python, BeautifulSoup, REST APIs) for batdongsan.com and 7+ other sites, plus the company's CSV/Excel/Sheets file sources via **Airbyte** – loading 300K+ property records into PostgreSQL.
- **Designed a 15-table dimensional data warehouse** (star schema) as a single source of truth for customers, properties and interactions, orchestrated end-to-end with **Airflow**.
- **Built a low-code CRM on AppSheet** (customer tracking, attendance, call logging for customers/landlords, apartment images, care history) – **cutting customer-care and property-lookup time by 50%**.
- **Delivered 10+ Looker dashboards** for customer demand, inventory and sales performance, used by 25+ users.

Data Research Assistant | AISIA Research Lab **Jun 2022 – Jan 2023**

- Labeled 80+ hours of Vietnamese speech data and helped optimize **ASR models (Wav2Vec 2.0)**, cutting word error rate by ~**12%** to narrow the gap with FPT's solution.

Personal Projects

Mini Lakehouse (Apache Iceberg Lab) **2026**
Technologies: Apache Iceberg, Spark (PySpark), Trino, Polaris, MinIO, Parquet, Docker repo

- Built a multi-engine lakehouse where Spark and Trino share Apache Iceberg tables via a Polaris REST catalog over MinIO, covering partitioning, time-travel, COW/MOR and table maintenance end-to-end.

FB Feed Crawler (Multi-Provider Scraping Platform) **2026**
Technologies: Python, FastAPI, PostgreSQL, APScheduler, Docker, Playwright/Crawl4AI repo

- Built a resilient Facebook crawler with a 3-engine fallback chain (Crawl4AI, Scrapling, facebook-scraper) behind circuit breakers, plus daily scheduling, crash recovery and a FastAPI management API.

CryptoStream Labs (Automated ELT Pipeline) **2024**
Technologies: GCP (GKE, BigQuery), Apache Airflow, dbt, Airbyte, Docker, Helm, SQL repo

- Built an ELT pipeline ingesting crypto data via Airbyte into BigQuery, transformed with dbt (partitioned OBT), orchestrated on Airflow and deployed on GKE via Helm.

Sign2Text (Traffic-Sign Recognition – University Research) **2022**
Technologies: Python, YOLOv7, VietOCR, OpenCV, Streamlit repo

- Led a university scientific-research team building a traffic-sign recognition system – labeled the dataset, trained a YOLOv7 detector, and extracted Vietnamese text from signs with VietOCR.

Education

University of Information Technology, VNU-HCM **2019 – Jun 2023**
Bachelor of Data Science Ho Chi Minh City, VN
GPA: 7.9/10 | TOEIC 725